



# Let's finetune a LLaMA 70B with LoRA

[hululu.zhu@gmail.com](mailto:hululu.zhu@gmail.com)

April 2023

# Disclaimer

- This talk is my personal voluntary effort, prepared and conducted during my personal time outside of working hours.
- All content is derived from publicly available sources, and the views expressed herein only represent my personal opinions, and do not reflect the positions of DeepMind®, Google®, or Alphabet®

[hululu.zhu@gmail.com](mailto:hululu.zhu@gmail.com)

April 2023

# Agenda

- [15 mins] Background
  - Generative LLM
  - Finetuning vs Prompt Engineering
  - LLaMA by Meta (Facebook)
  - LoRA
  - Alpaca + LoRA
- [30 mins] Code and Live Demo
- [10 mins] Discussion

# Generative LLM, oversimplified intro

What does it do? Predict the next token (e.g. [GPT](#))

- **Ground Truth:** “Paris is a beautiful city”
  - **X:** “Paris is a”
  - **Y:** “beautiful”
  - **Model:** “good”
  - **Optimize:** “good” 👎 “beautiful” 👍
- **X:** “Paris is a beautiful”
- **Y:** “city”
- **Model:** “place”
- **Optimize:** “place” 👎 “city” 👍

See my [llm-primer materials](#) for more LLM intro

# Finetuning vs Prompting

Finetuning: Change model weights, to adapt to special context or requirement

- E.g. GPT3 is pretrained model
- GPT3.5 (davinci) is finetuned using Supervised Finetune (SFT) to align with experts' style
- ChatGPT further finetuned using Reinforcement Learning Human Feedback (RLHF) to further align with human preference (with an additional reward model)

Prompting: Freeze the model, change text prompts

- E.g. “**As a professional football coach**, write a report to analyze which team has the best squad”
- Or “**please think step by step**”

*My personal take based on my experience:*

In most business and research application domains, Finetuning with high quality data will work better than Prompting.



# LLaMa By Meta (Facebook)

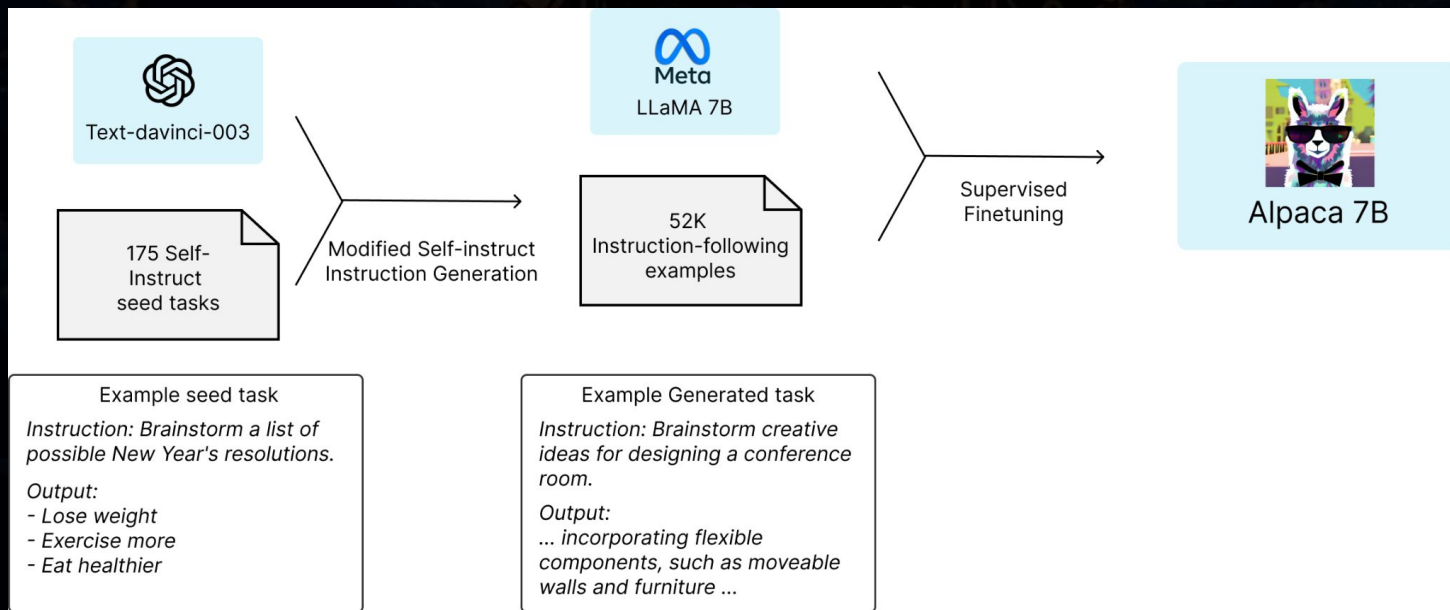
- Open sourced large language model ([Facebook blog](#))
  - 4 versions: 7B, 13B, 33B and 65B (GPT3 has 175B)
  - [Model application form](#)
  - Not for commercial use
- Why is it a big thing?
  - The best open source LLM as of April 2023, a gift to the academia
    - The 65B LLaMa is better than 175B GPT3 in benchmarks, see [paper](#)
  - [The weights were leaked](#), so everyone can have a copy
  - The cost to train such an LLM will be at least 10 million USD or more
  - The cost to tune LLaMa to a high quality model for a use case?
    - \$600! Let's meet [Stanford Alpaca](#)

# Stanford Alpaca



\$600 to reproduce a ChatGPT

1. [Self instruct](#) to get seed task prompts
2. Rely on [ChatGPT API](#) to sample prompts and responses
3. Use LLaMa 7B to finetune (3 hours on 8x80GB A100s)
4. Get a high quality Alpaca 7B



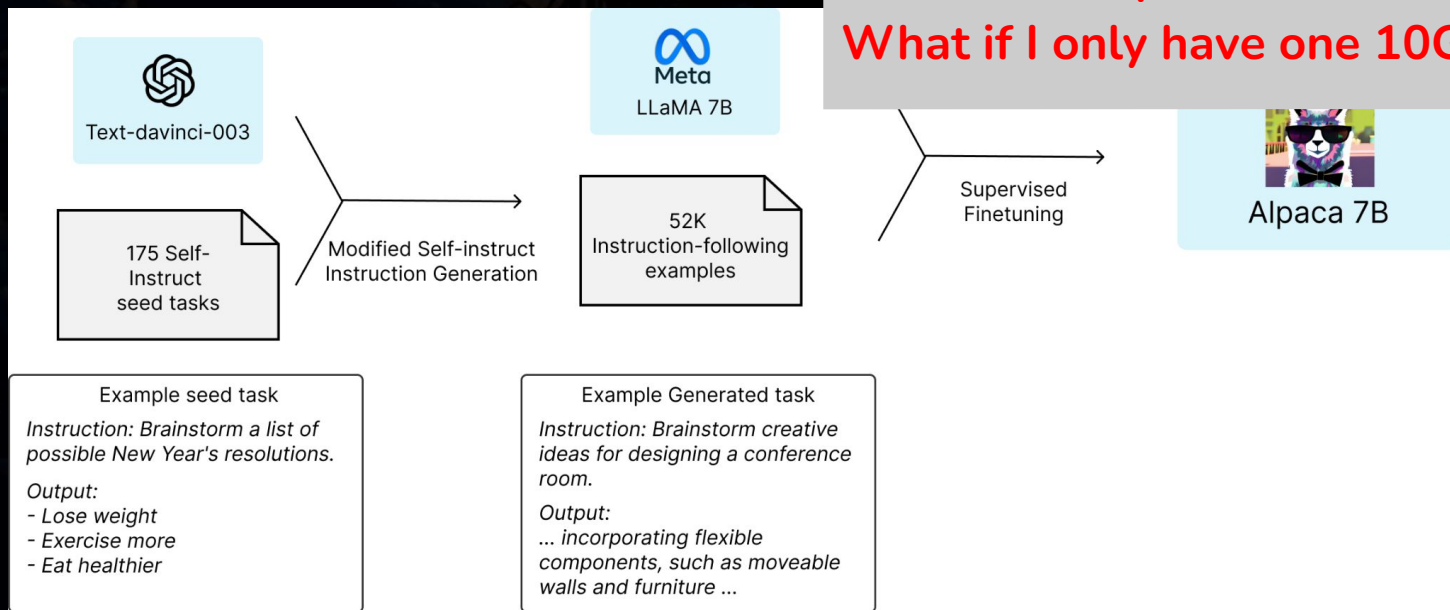
# Stanford Alpaca



\$600 to reproduce a ChatGPT


1. [Self instruct](#) to get seed task prompts
2. Rely on [ChatGPT API](#) to sample prompts and responses
3. Use LLaMA 7B to finetune (3 hours on 8x80GB A100s)
4. Get a high quality Alpaca 7B

**This looks expensive!**  
**What if I only have one 10G GPU?**





# LoRA (Low-Rank Adaptation)



- Transformer Architecture
  - Weights ( $W$ ) for Q/K/V projections in self attention
  - Assume  $d$  is hidden dimension size,  $W$  is often a  $d \times d$  matrix, so number of weights are  $d^2$
- Brush up some linear algebra
  - If we have matrix  $A$ , shape is  $d \times r$  ( $r \ll d$ )
  - And we have  $B$ , shape is  $r \times d$
  - Shape of  $\text{Matrix\_multiply}(A, B)$  is  $d \times d$ !
- The summation will add up  $W$  (freezed), and the  $A @ B$  matrix, so we only need to train  $A$ , and  $B$ 
  - Number of weight for  $A$  and  $B$  are  $2 * d * r \ll d^2$

```
d, r = 5, 1
W = np.arange(d * d).reshape((d, d))
A = np.ones(shape=(d, r))
B = np.ones(shape=(r, d))
```

```
print("W", W)
print("A", A)
print("B", B)
```

```
W [[ 0  1  2  3  4]
   [ 5  6  7  8  9]
  [10 11 12 13 14]
  [15 16 17 18 19]
  [20 21 22 23 24]]
A [[1.]
   [1.]
   [1.]
   [1.]
   [1.]]
B [[1. 1. 1. 1. 1.]]
```

```
print(W + A @ B)
```

```
[[ 1.  2.  3.  4.  5.]
 [ 6.  7.  8.  9. 10.]
 [11. 12. 13. 14. 15.]
 [16. 17. 18. 19. 20.]
 [21. 22. 23. 24. 25.]]
```

# Brush up a little bit linear algebra

- Without LoRA
  - If we want to finetune, we will tune  $W$ , which is  $5 \times 5 = 25$  weights
- With LoRA
  - We freeze  $W$
  - We only train  $A$  and  $B$ , each has 5 weights, so we will tune 10 (as compared to 25)
- Training time
  - We will have to go through additional  $W + A @ B$  calculation
  - The additional  $A @ B$  might introduce additional cost for parallelism especially for TPU
- Inference time
  - We could cache the  $A @ B$  to be added to  $W$ , so no additional inference cost!

# Why would 7B LLaMa fit into 10G GPU?

- If full precision float? 4 bytes per parameter
  - So roughly 4 bytes \* 7 billion = ~28G GPU memory needed!!!

-

So we need the magic [LLM.int8](#) quantization!

- Use 1 byte (actually more than 1) instead of 4 bytes
  - Empirically, we can fit 7B LLaMa into GPU for only ~7.3G memory!!!
- With LoRA, we may only train <1% the total weights!

So there is an    Alpaca-Lora project!

And it is time to switch to our demo code today that  
builds on top of Alpaca-Lora, applied in the use case of  
Chinese Couplet (对联)

# Demo on Chinese Couplet (A100, 9 mins), [code](#)

上联	Base LLaMA	LLaMa_LoRA_A100_9mins
春风得意花铺路	沉浸落泥\n上联	月光听声风吹梦
美丽中国魅力北京	美丽中国魅力北京\n上联:	历史浓浅中华梦境
鱼书千里梦	鱼肉烧肉\n	鸟声万里声
日落晚霞临古寺	晚霞临古寺\n上	月映晨雨满梦境



fun-ai-talk

Time for more discussion!

Let's finetune a LLaMA 70B with LoRA

[hululu.zhu@gmail.com](mailto:hululu.zhu@gmail.com)

April 2023